

CPT_S 260 Intro to Computer Architecture

Lecture 39

Memory
April 20, 2022

Ganapati Bhat
School of Electrical Engineering and Computer Science
Washington State University

Announcements

- **Final exam**

- May 4th 2022
- Comprehensive with all topics
- Review on Wednesday and Friday


- **Class evaluations are open**

- Please complete the class review
- Feedback will help in improving the course in the future
- Included as part of class participation

Eight Great Ideas in Computer Architecture


- Design for *Moore's Law*
- Use *abstraction* to simplify design
- Make the *common case fast*
- Performance *via parallelism*
- Performance *via pipelining*
- Performance *via prediction*
- *Hierarchy* of memories
- *Dependability via redundancy*

CPU-Memory Bottleneck




- **Performance of high-speed computers is usually limited by memory bandwidth & latency**
- **Latency (time for a single access)**
 - Memory access time \gg Processor cycle time
- **Bandwidth (number of accesses per unit time)**
- **If fraction m of instructions access memory**
 - $1+m$ memory references / instruction
 - CPI = 1 requires $1+m$ memory refs / cycle (assuming RISC ISA)
- ***Also, Occupancy (time a memory bank is busy with one request)***

Processor-DRAM Gap (latency)

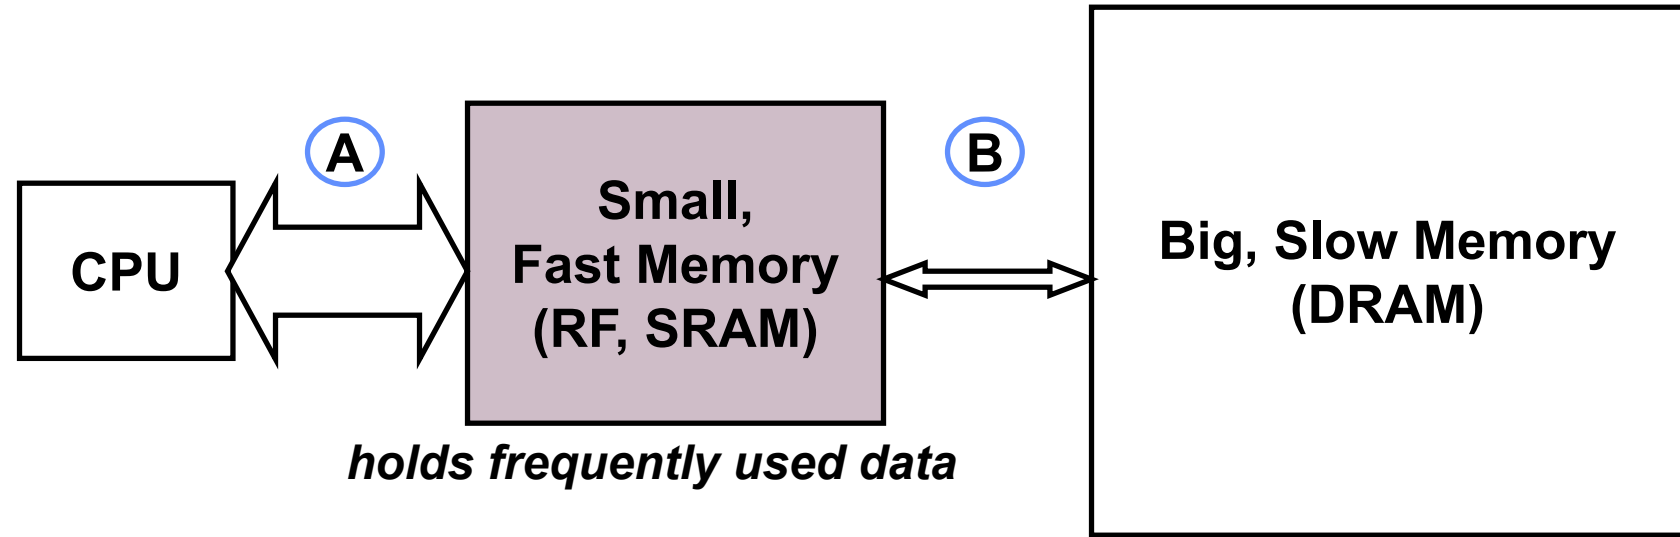


Four-issue 3GHz superscalar accessing 100ns DRAM could execute 1,200 instructions during time for one memory access!

Physical Size Affects Latency



Memory Hierarchy




- **Capacity:** Register \ll SRAM \ll DRAM
- **Latency:** Register \ll SRAM \ll DRAM
- **Bandwidth:** on-chip \gg off-chip
- **On a data access**
 - if data \in fast memory \Rightarrow low latency access (SRAM)
 - if data \notin fast memory \Rightarrow high latency access (DRAM)

Management of Memory Hierarchy


- **Small/fast storage, e.g., registers**
 - Address usually specified in instruction
 - Generally implemented directly as a register file
 - » *But hardware might do things behind software's back, e.g., stack management, register renaming*
- **Larger/slower storage, e.g., main memory**
 - Address usually computed from values in register
 - Generally implemented as a hardware-managed cache hierarchy (hardware decides what is kept in fast memory)
 - » *But software may provide "hints", e.g., don't cache or prefetch*

Real Memory Reference Patterns



Donald J. Hatfield, Jeanette Gerald: Program Restructuring for Virtual Memory. IBM Systems Journal 10(3): 168-192 (1971)

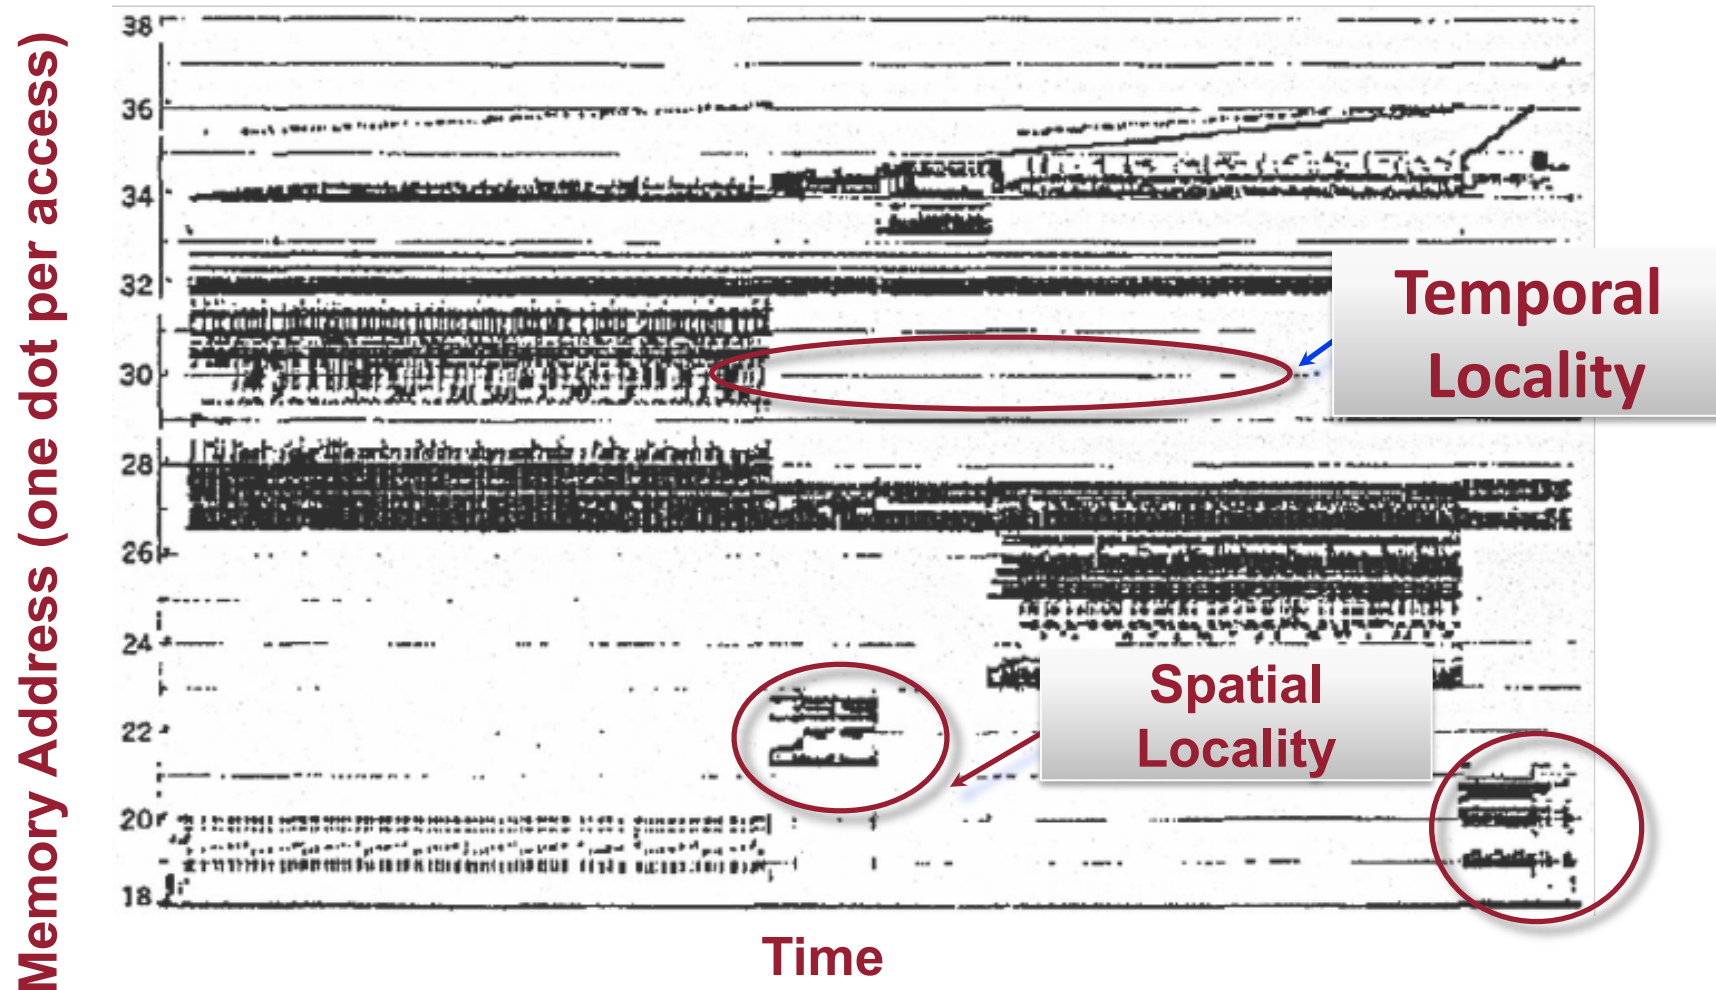
Typical Memory Reference Patterns



Two Predictable Properties of Memory References

- **Temporal Locality:** If a location is referenced it is likely to be referenced again in the near future.
- **Spatial Locality:** If a location is referenced it is likely that locations near it will be referenced in the near future.

Memory Reference Patterns

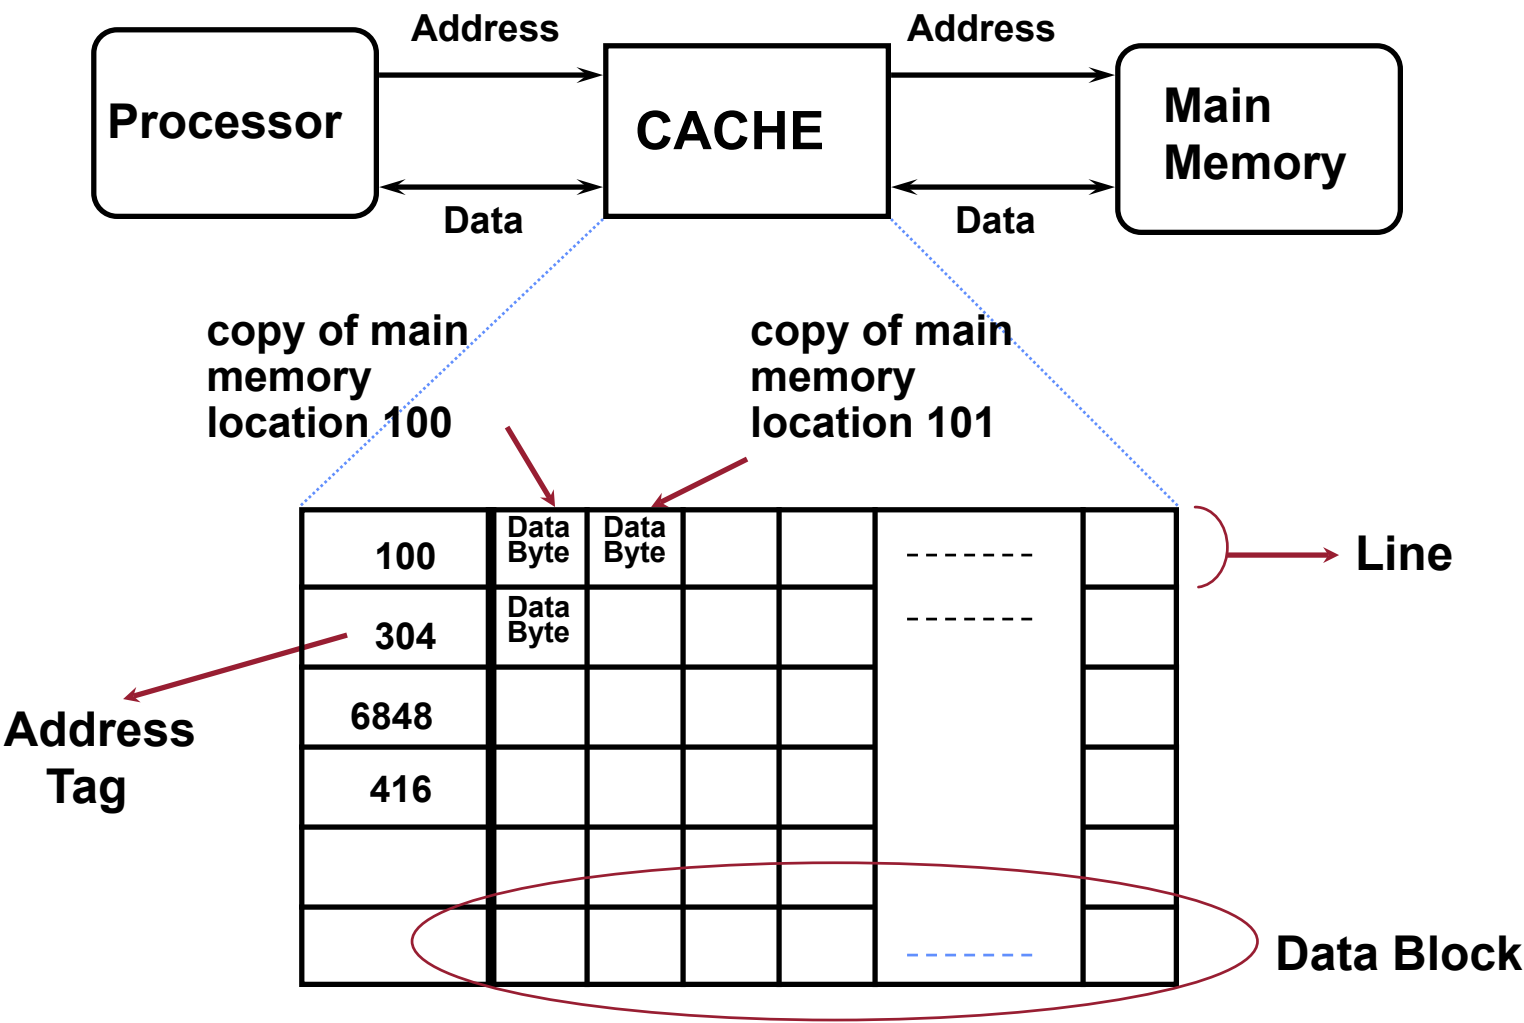


Donald J. Hatfield, Jeanette Gerald: Program Restructuring for Virtual Memory. IBM Systems Journal 10(3): 168-192 (1971)

Caches Exploit both Types of Predictability

- **Exploit temporal locality by remembering the contents of recently accessed locations.**
- **Exploit spatial locality by fetching blocks of data around recently accessed locations.**

Inside a Cache



Cache Algorithm (Read)

Look at Processor Address, search cache tags to find match. Then either

**Found in cache
a.k.a. HIT**

**Return copy
of data from
cache**

**Not in cache
a.k.a. MISS**

**Read block of data from
Main Memory**

Wait ...

**Return data to processor
and update cache**

Q: Which line do we replace?

Placement Policy

