

CPT_S 260 Intro to Computer Architecture

Lecture 35

Pipeline Datapath

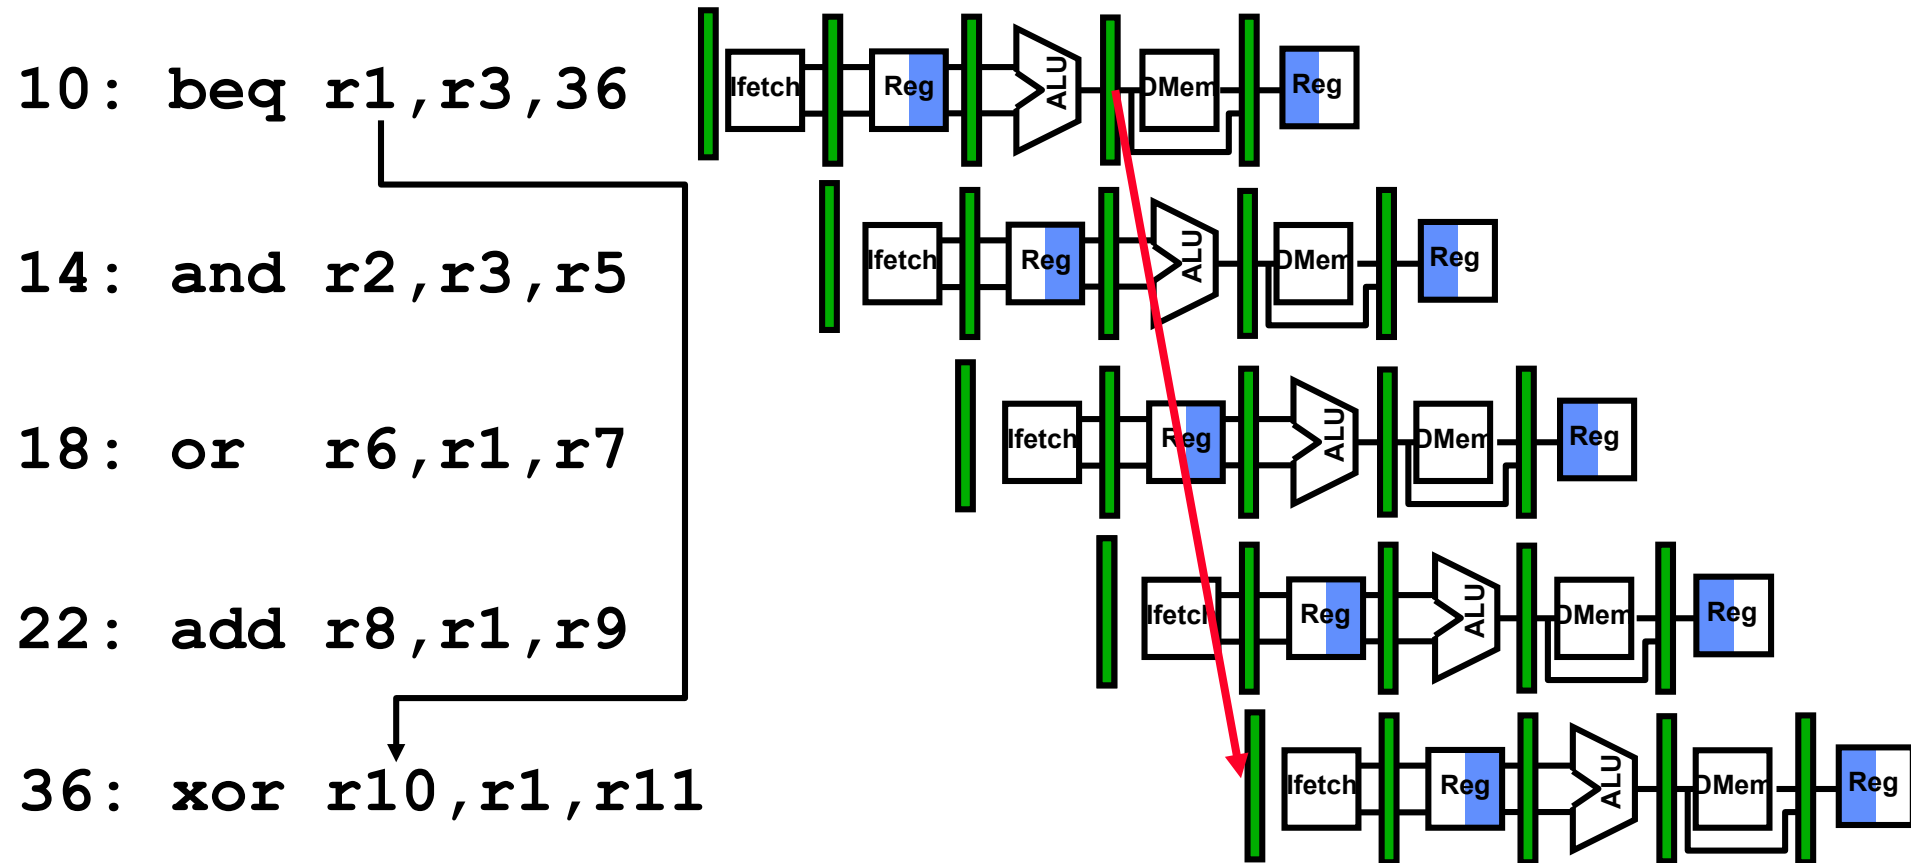
April 11, 2021

Ganapati Bhat
School of Electrical Engineering and Computer Science
Washington State University

Control Hazards

- **Branch determines flow of control**
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - » Still working on ID stage of branch
- **In MIPS pipeline**
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage

Control Hazard on Branches Three Stage Stall



What do you do with the 3 instructions in between?

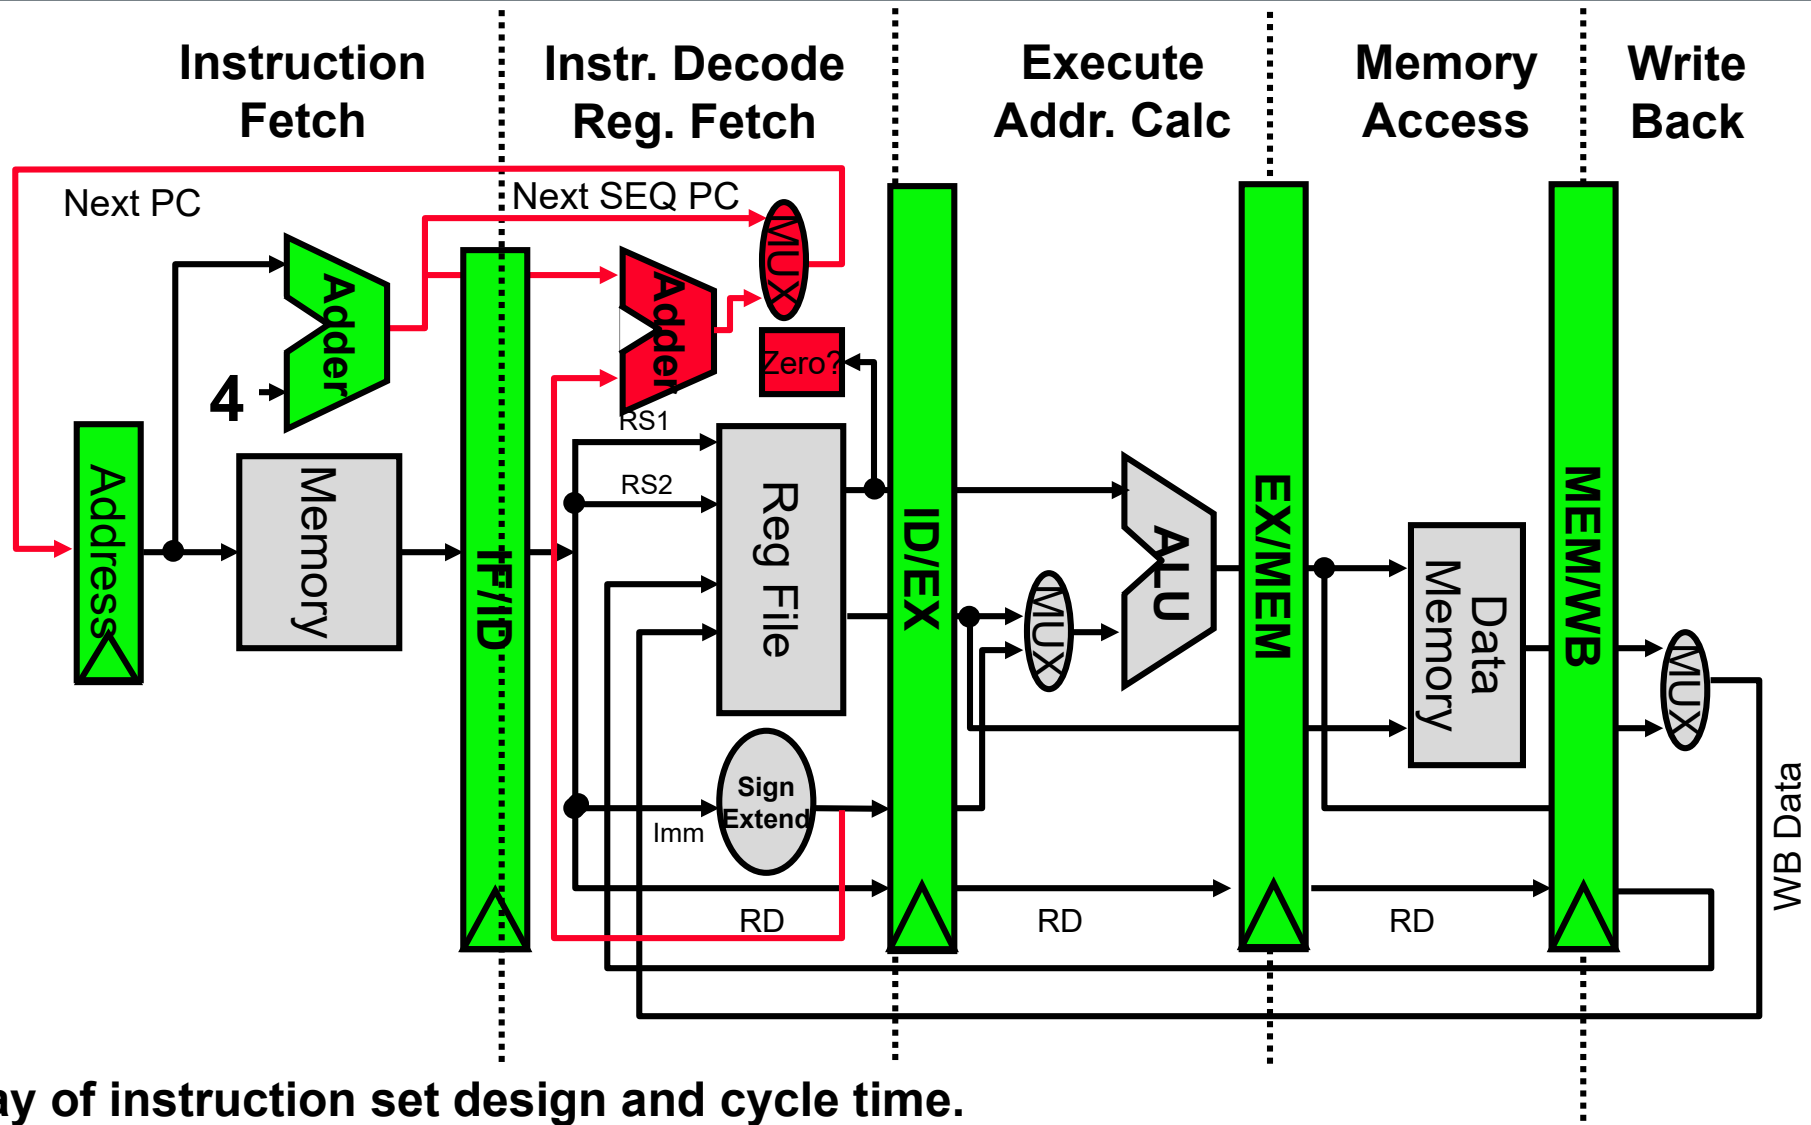
How do you do it?

Where is the “commit”?

Branch Stall Impact

- **If CPI = 1, 30% branch,
Stall 3 cycles => new CPI = 1.9!**
- **Two-part solution:**
 - Determine branch taken or not sooner, AND
 - Compute taken branch address earlier
- **MIPS branch tests if register = 0 or \neq 0**
- **MIPS Solution:**
 - Move Zero test to ID/RF stage
 - Adder to calculate new PC in ID/RF stage
 - 1 clock cycle penalty for branch versus 3

Pipelined MIPS Datapath



Four Branch Hazard Alternatives

1: Stall until branch direction is clear

2: Predict Branch Not Taken

- Execute successor instructions in sequence
- “Squash” instructions in pipeline if branch actually taken
- Advantage of late pipeline state update
- 47% MIPS branches not taken on average
- PC+4 already calculated, so use it to get next instruction


3: Predict Branch Taken

- 53% MIPS branches taken on average
- But haven't calculated branch target address in MIPS
 - » MIPS still incurs 1 cycle branch penalty
 - » Other machines: branch target known before outcome
- What happens when hit not-taken branch?

Four Branch Hazard Alternatives

4: Delayed Branch

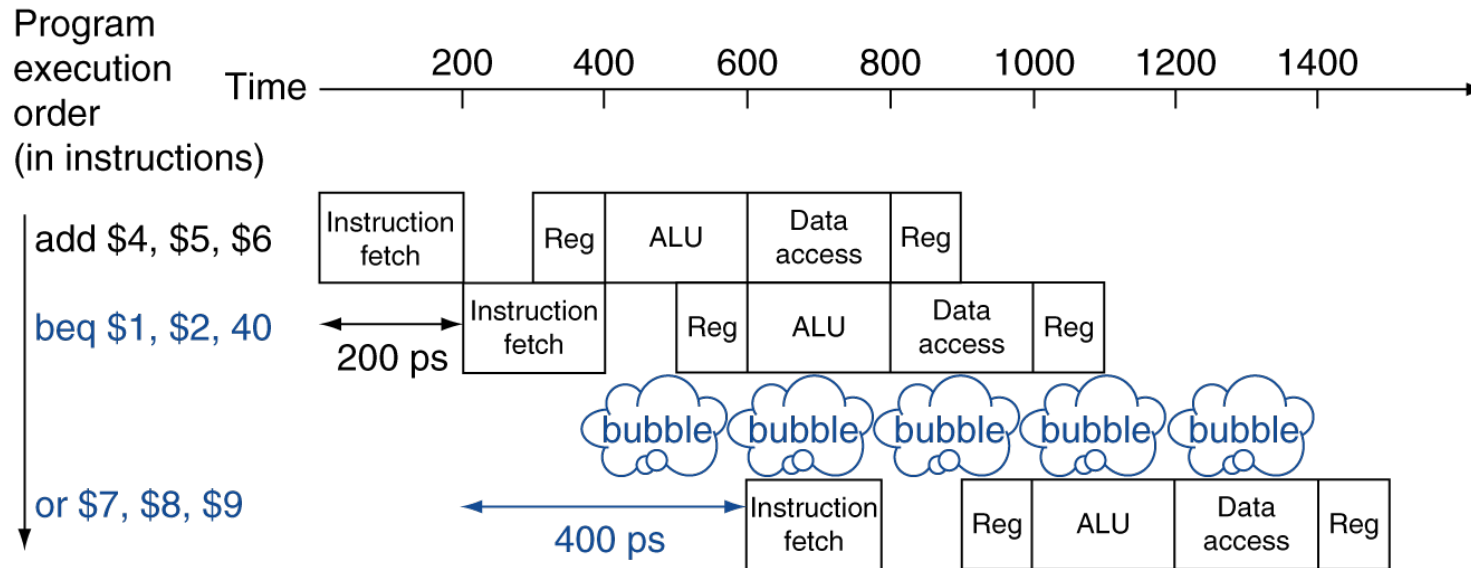
- Define branch to take place **AFTER** a following instruction



- 1 slot delay allows proper decision and branch target address in 5 stage pipeline
- MIPS uses this

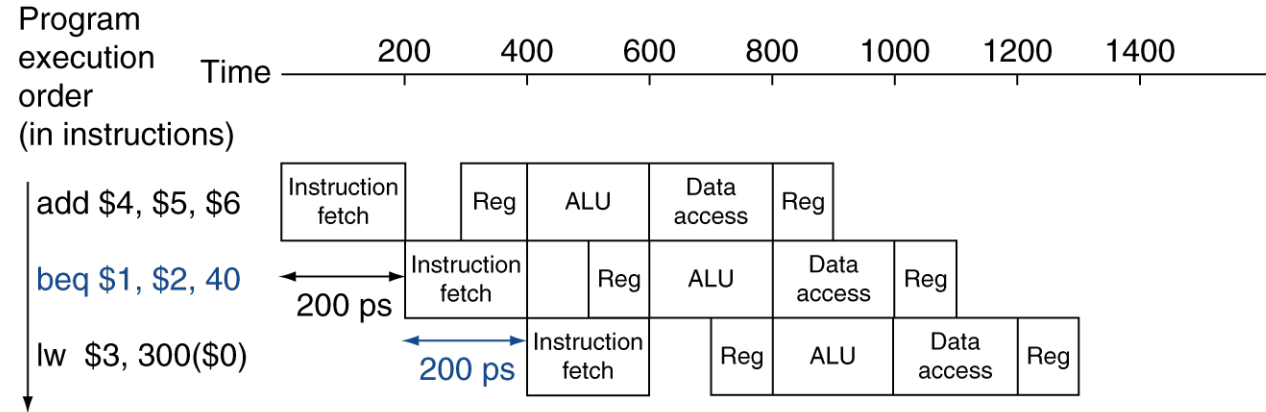
Stall on Branch

- Wait until branch outcome determined before fetching next instruction

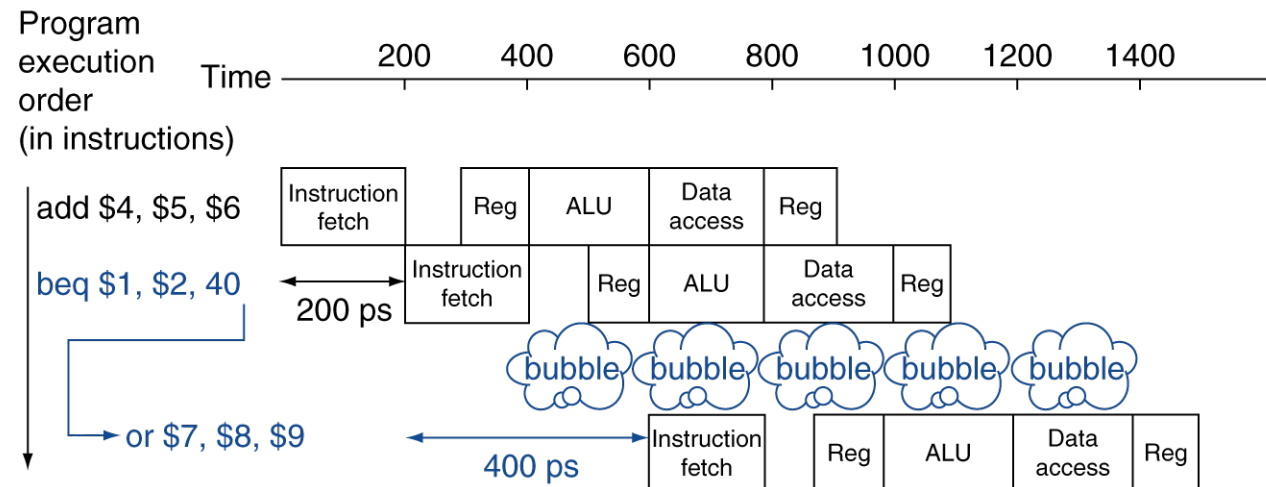


MIPS with Predict Not Taken

Prediction correct

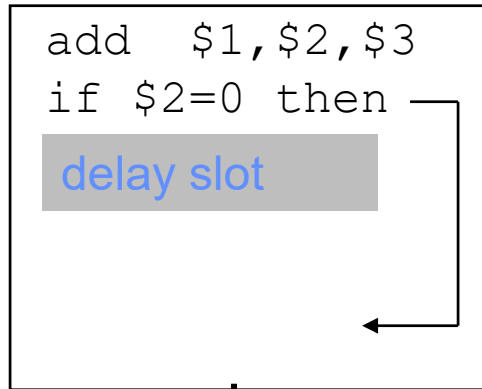


Prediction incorrect




Scheduling Branch Delay Slots

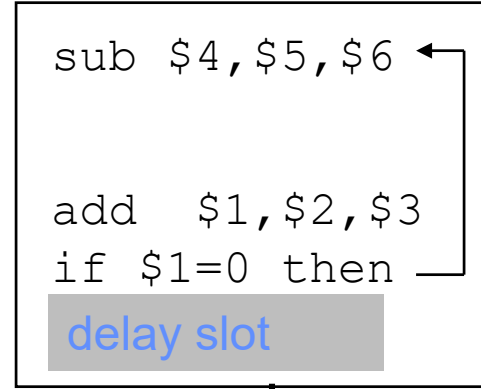
A. From before branch



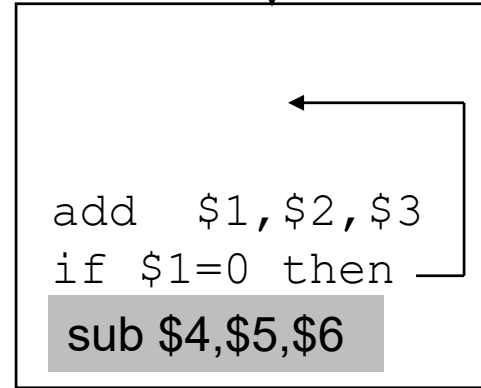
becomes



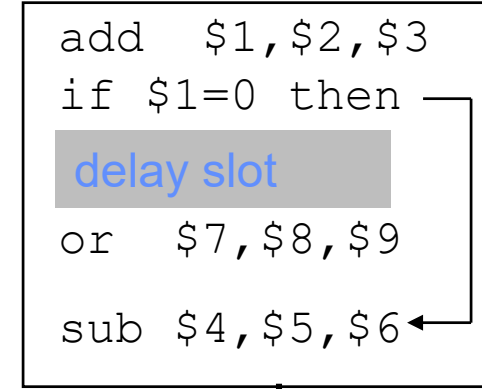
B. From branch target



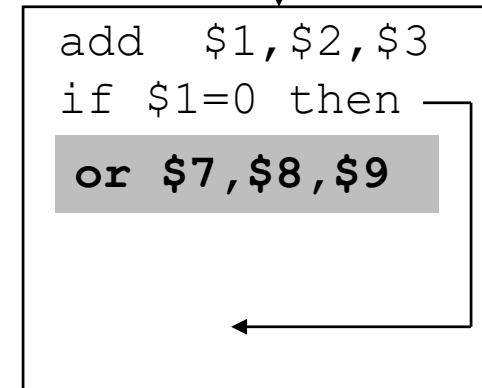
becomes



C. From fall through



becomes



- A is the best choice, fills delay slot & reduces instruction count (IC)
- In B, the `sub` instruction may need to be copied, increasing IC
- In B and C, must be okay to execute `sub/or` when branch fails

More-Realistic Branch Prediction

- **Static branch prediction**

- Based on typical branch behavior
- Example: loop and if-statement branches
 - » Predict backward branches taken
 - » Predict forward branches not taken

- **Dynamic branch prediction**


- Hardware measures actual branch behavior
 - » e.g., record recent history of each branch
- Assume future behavior will continue the trend
 - » When wrong, stall while re-fetching, and update history

Delayed Branch

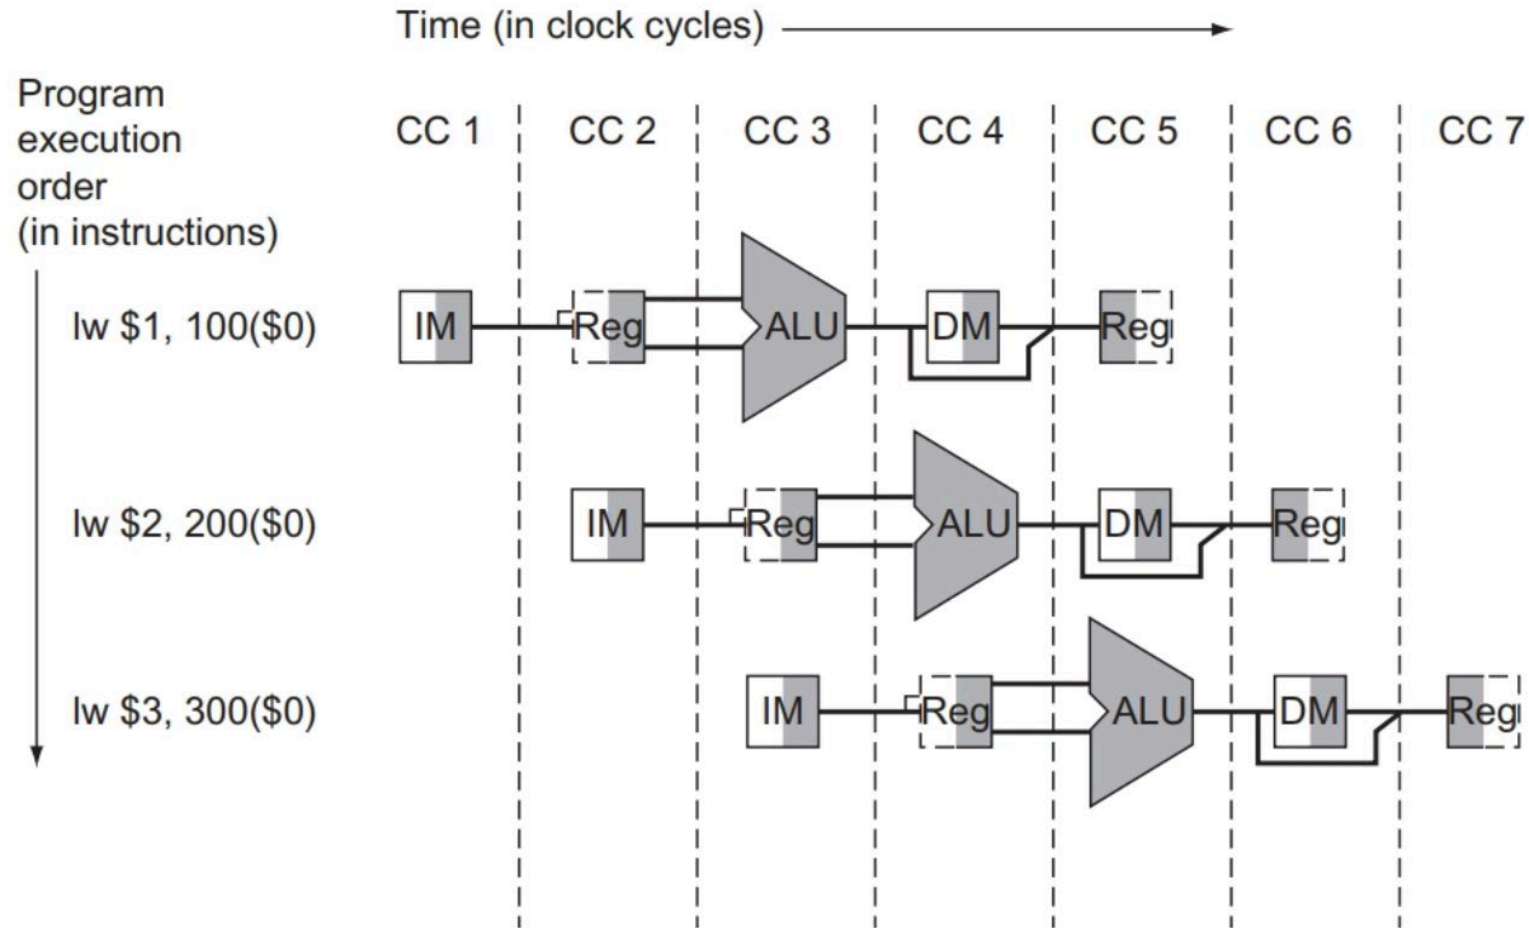
- **Compiler effectiveness for single branch delay slot:**
 - Fills about 60% of branch delay slots
 - About 80% of instructions executed in branch delay slots useful in computation
 - About 50% (60% x 80%) of slots usefully filled
- **Delayed Branch downside: As processor go to deeper pipelines and multiple issue, the branch delay grows and need more than one delay slot**
 - Delayed branching has lost popularity compared to more expensive but more flexible dynamic approaches
 - Growth in available transistors has made dynamic approaches relatively cheaper

Pipeline Datapath

MIPS Pipelined Datapath

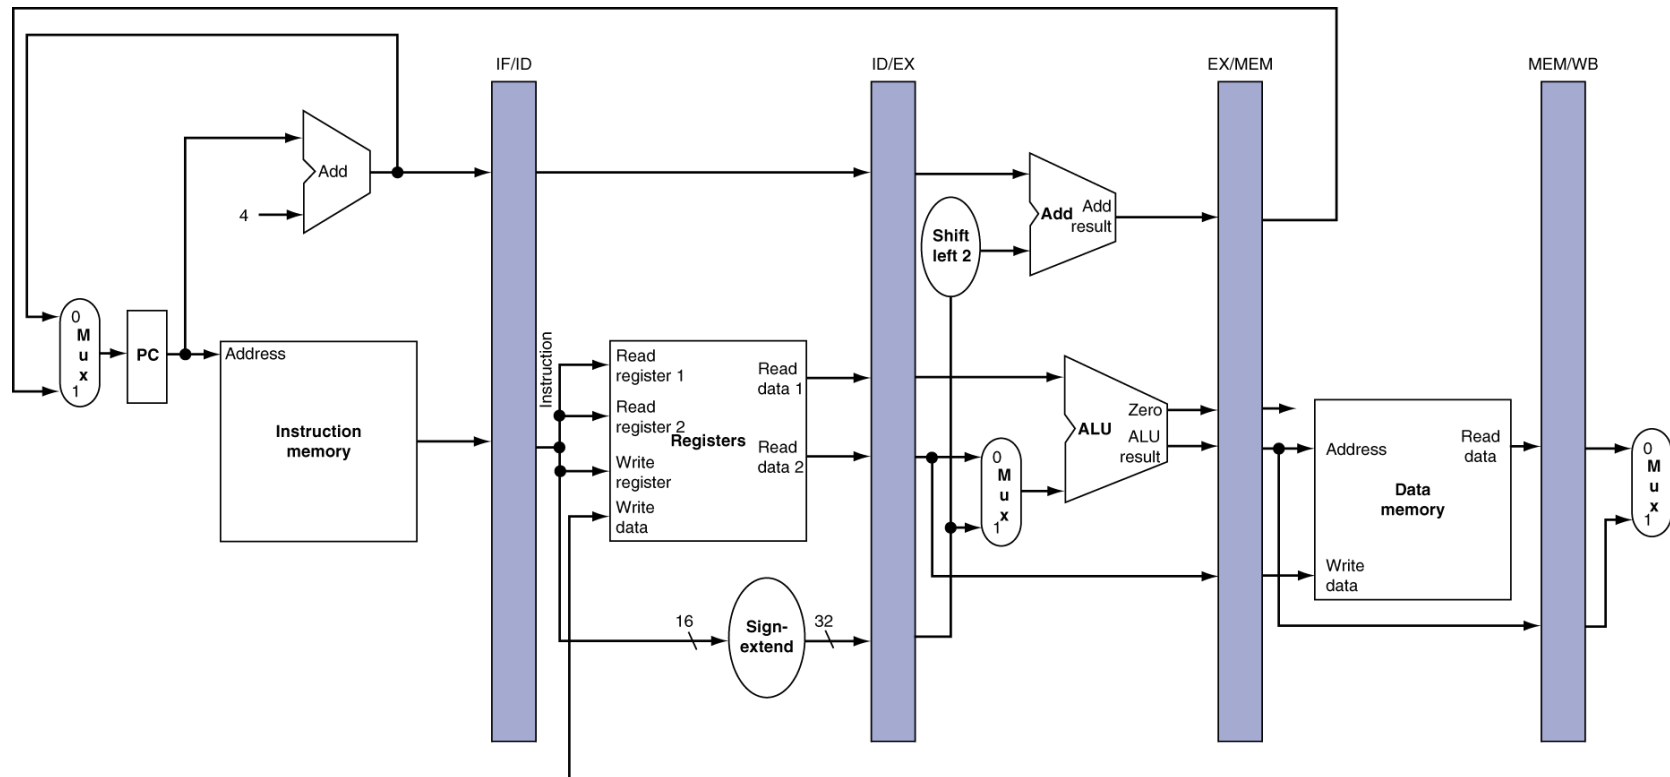


MIPS Pipelined Datapath



Pipeline registers

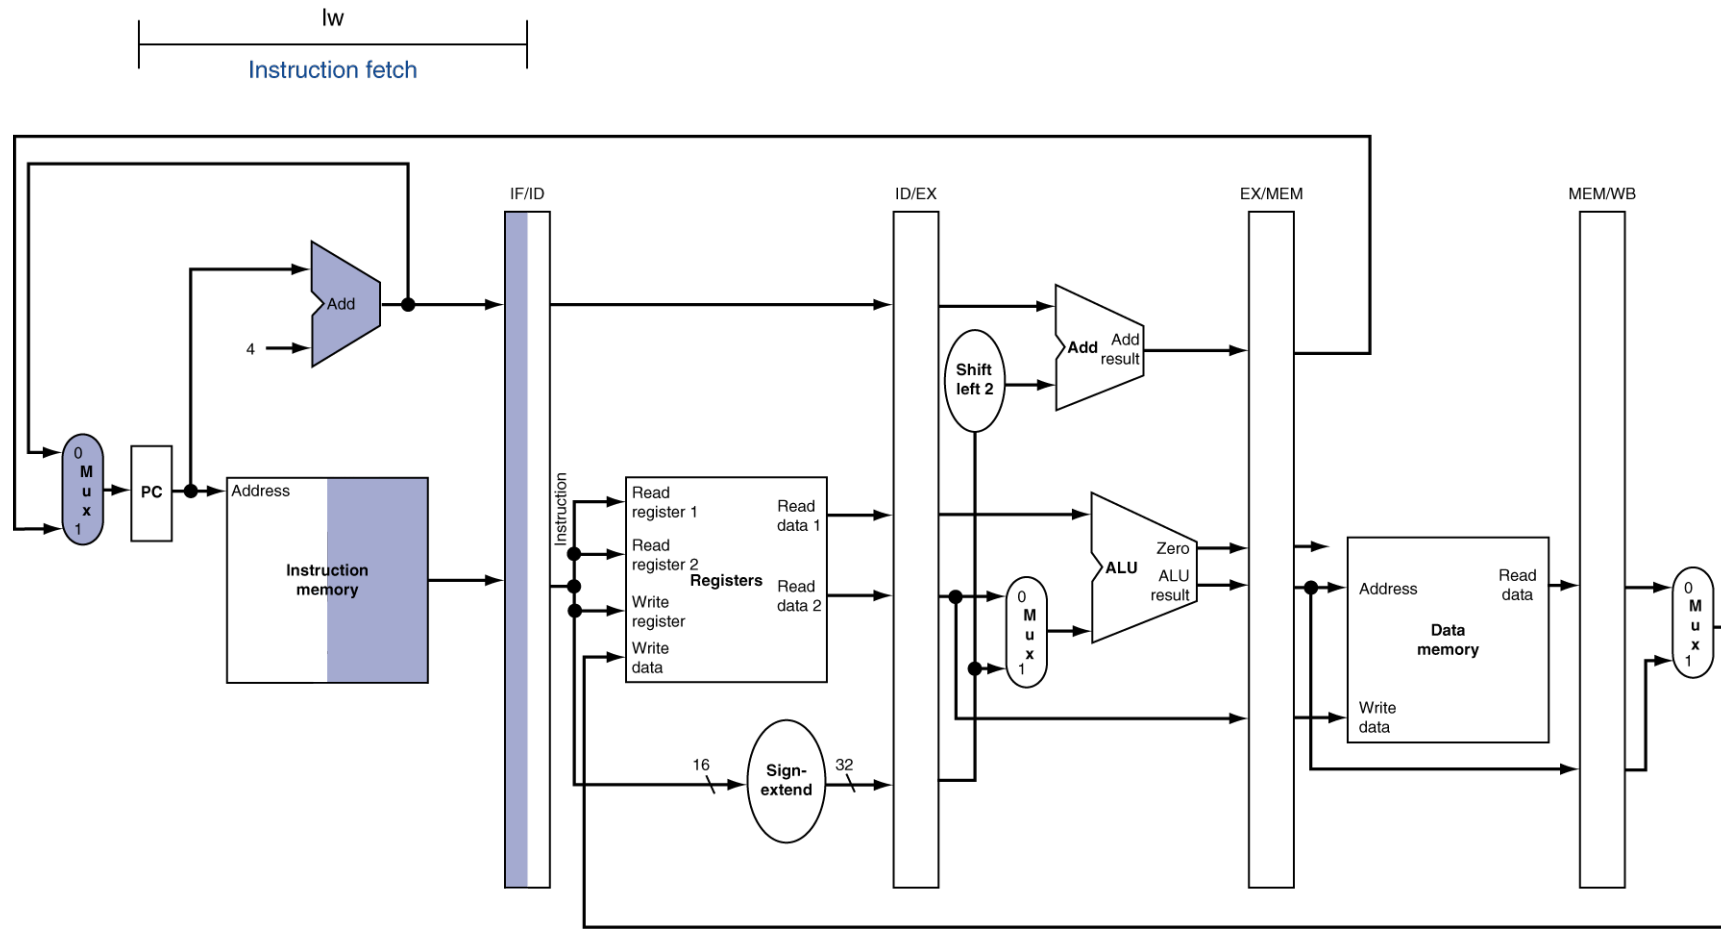
- **Need registers between stages**
 - To hold information produced in previous cycle




Pipeline Operation

- **Cycle-by-cycle flow of instructions through the pipelined datapath**
 - “Single-clock-cycle” pipeline diagram
 - » Shows pipeline usage in a single cycle
 - » Highlight resources used
 - c.f. “multi-clock-cycle” diagram
 - » Graph of operation over time
- **We’ll look at “single-clock-cycle” diagrams for load & store**


IF for Load, Store, ...



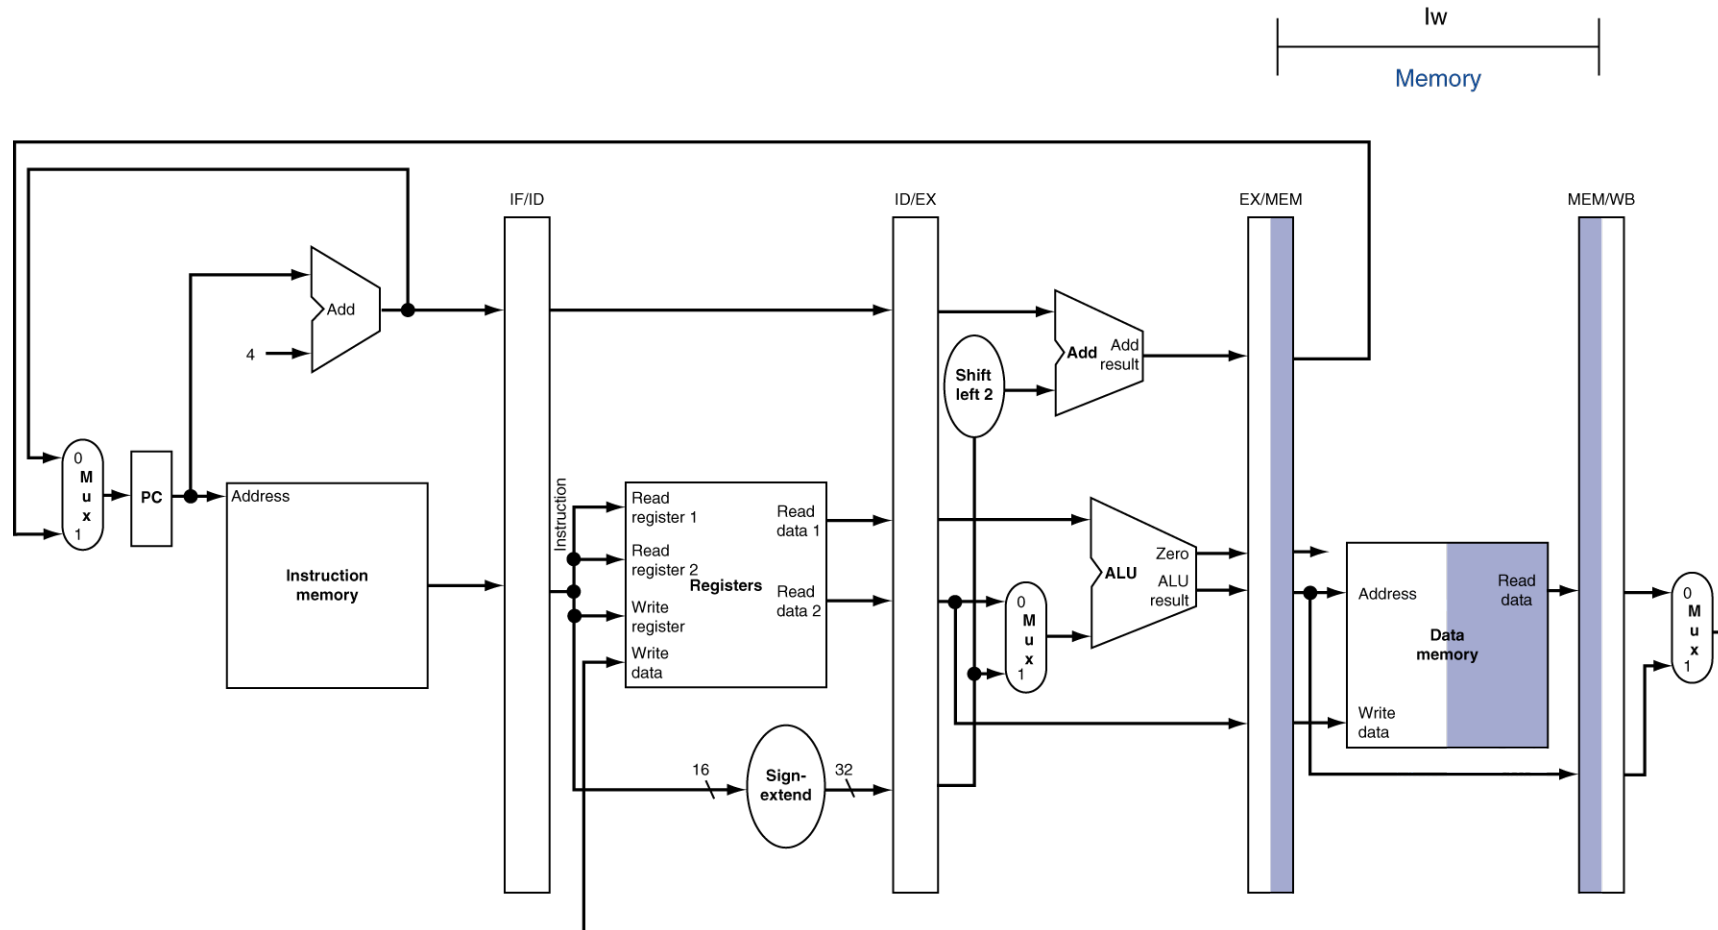
ID for Load, Store, ...



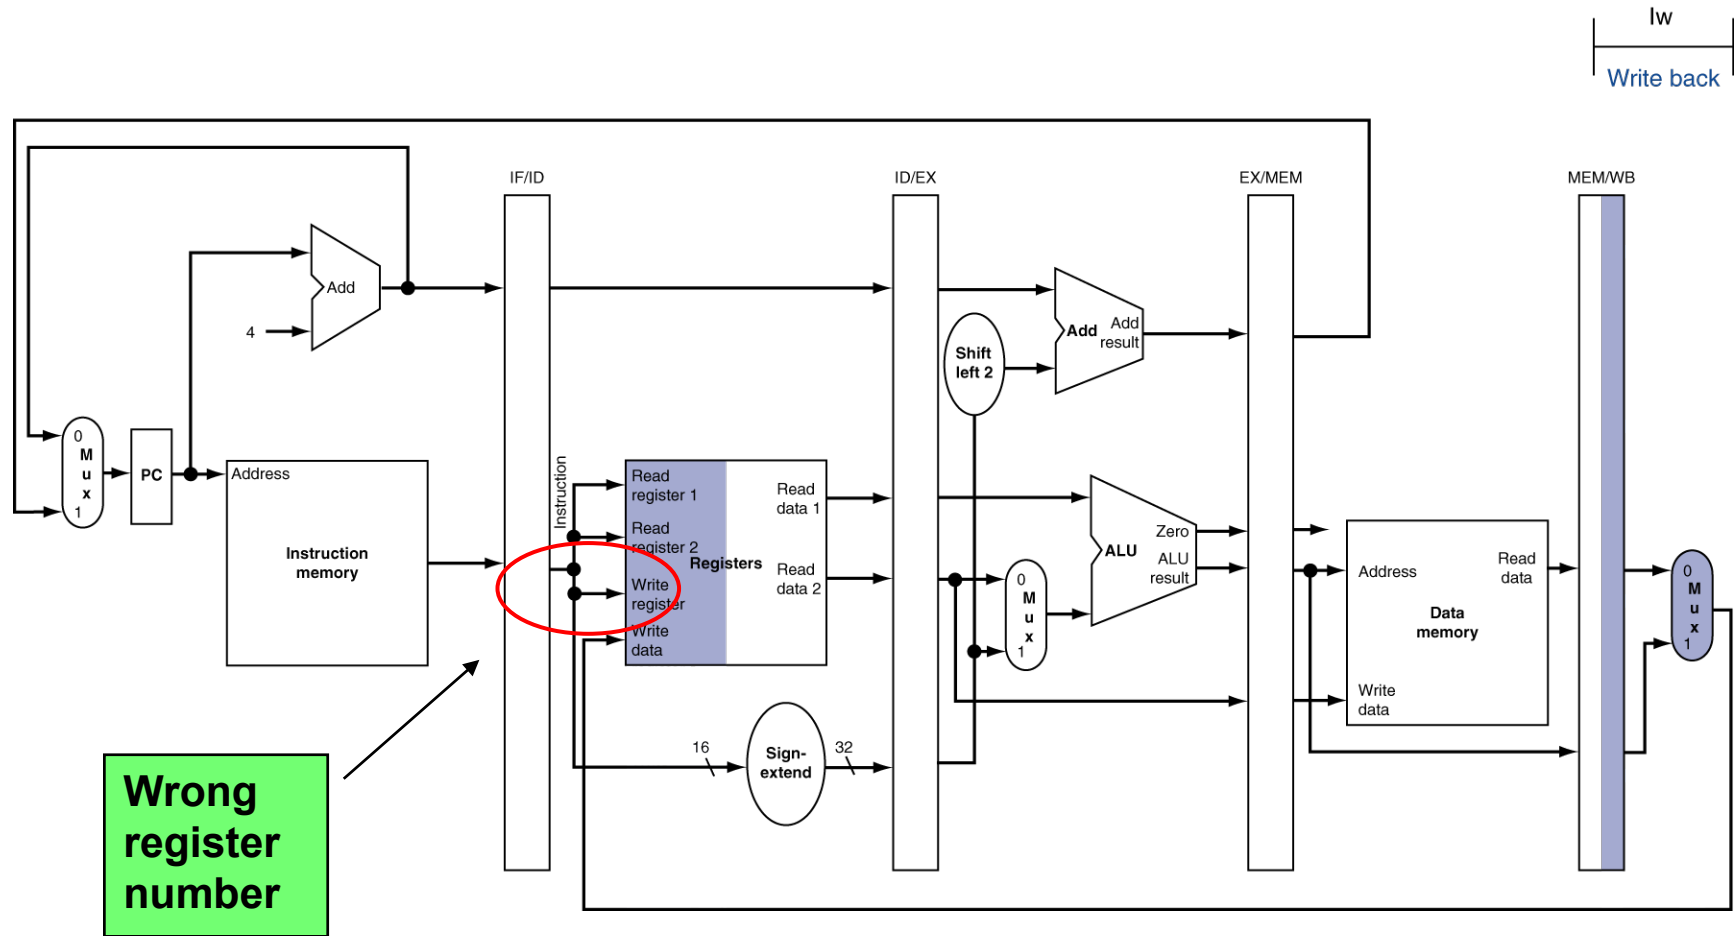
EX for Load



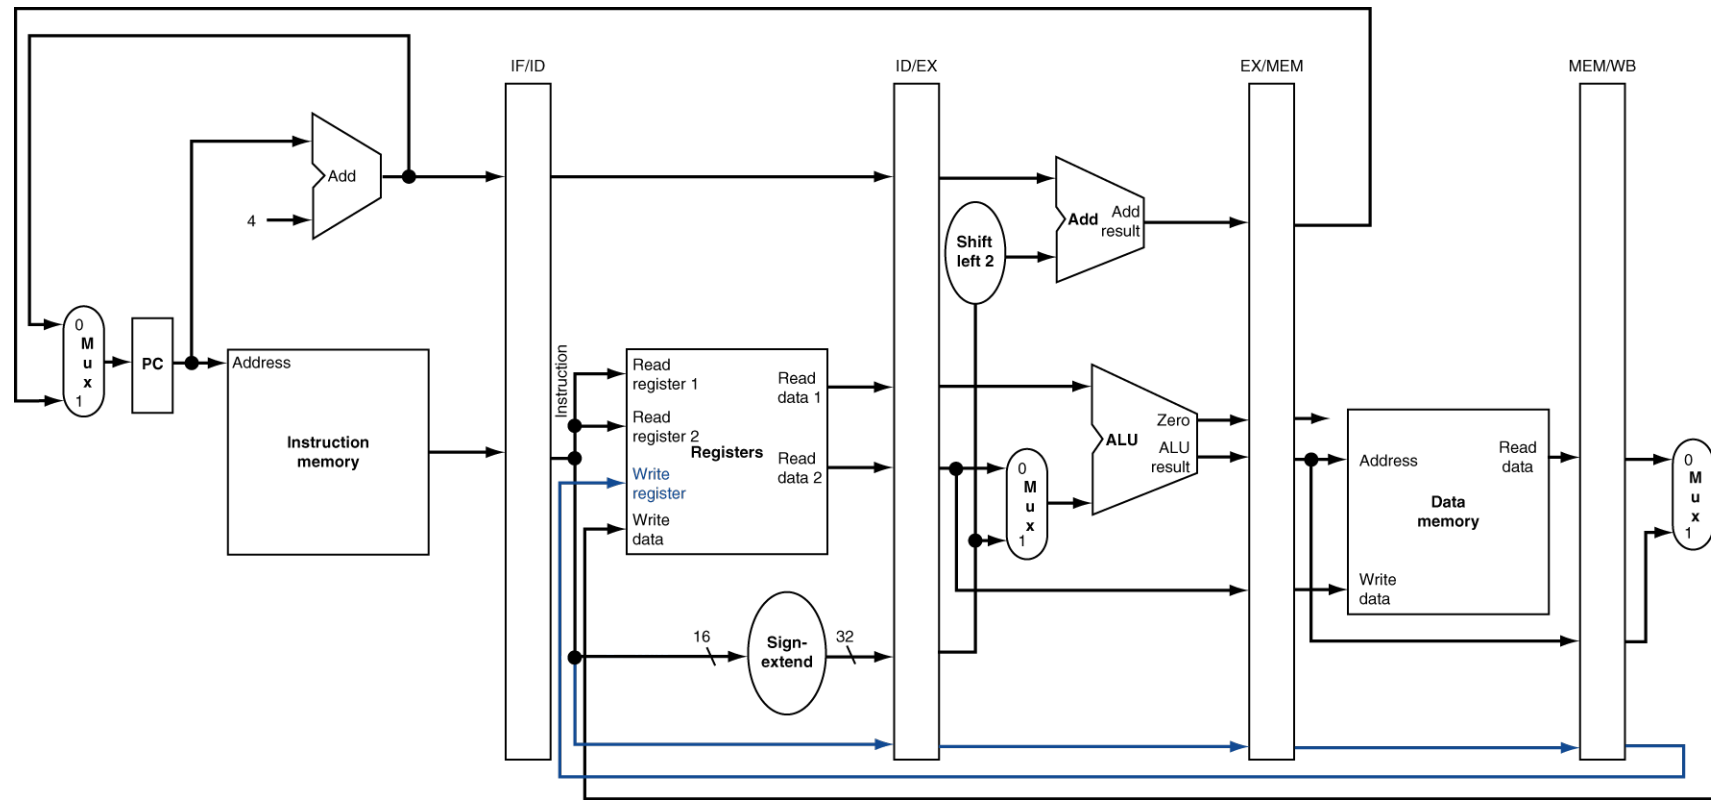
MEM for Load



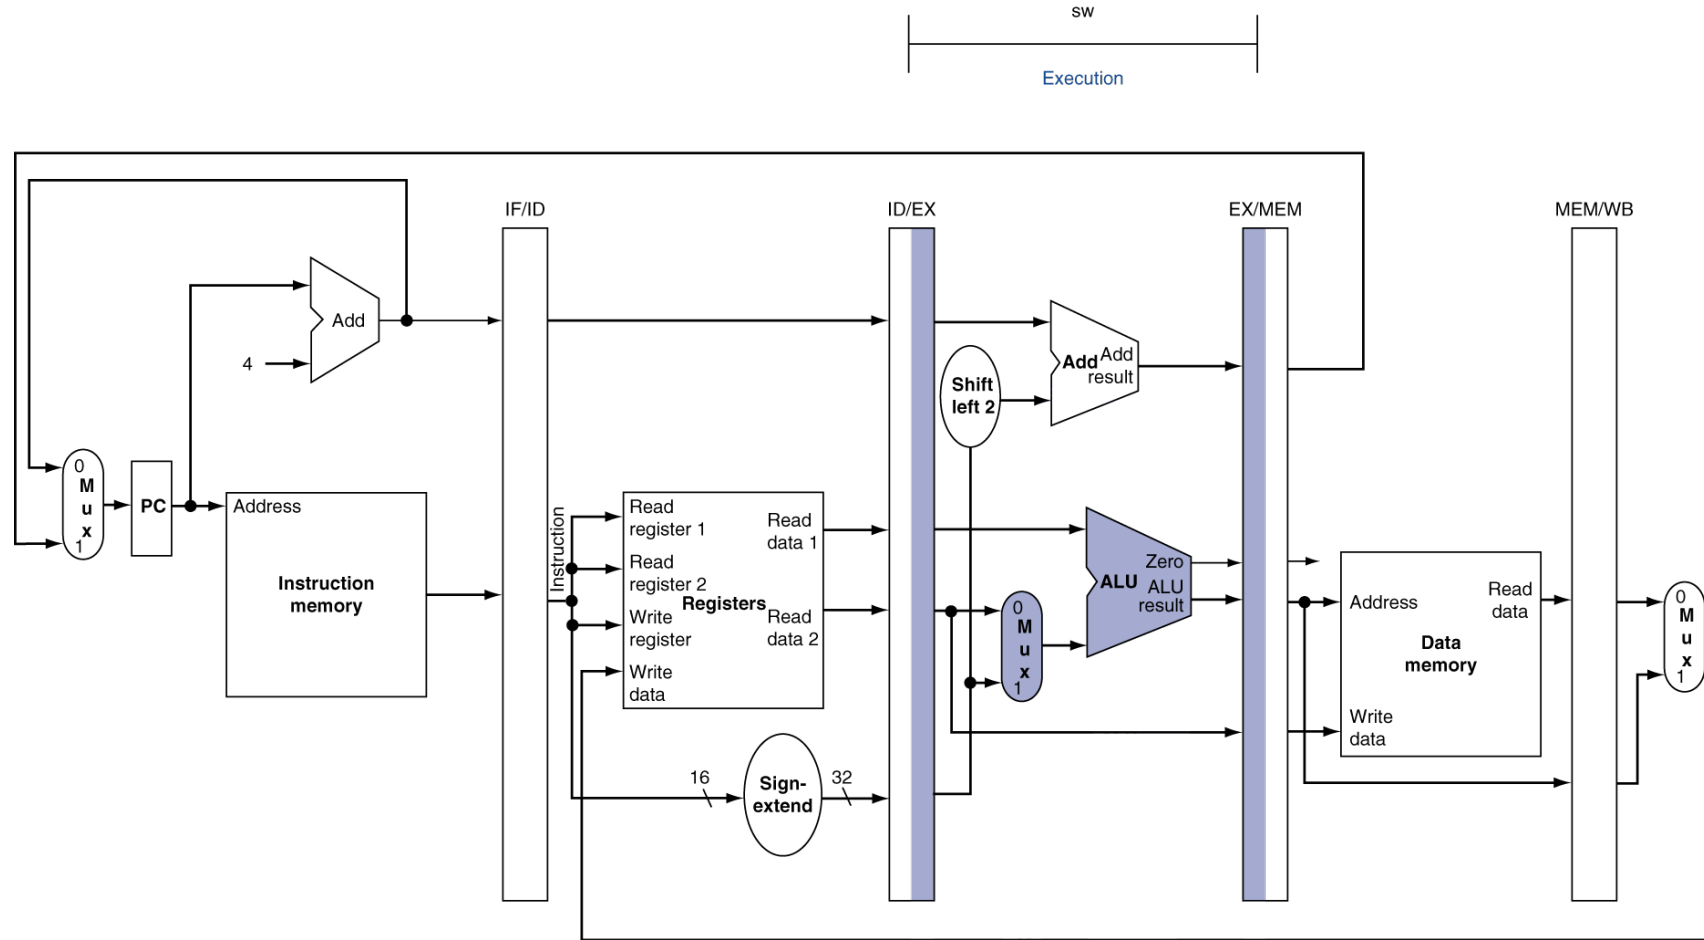
WB for Load




Corrected Datapath for Load




EX for Store



MEM for Store



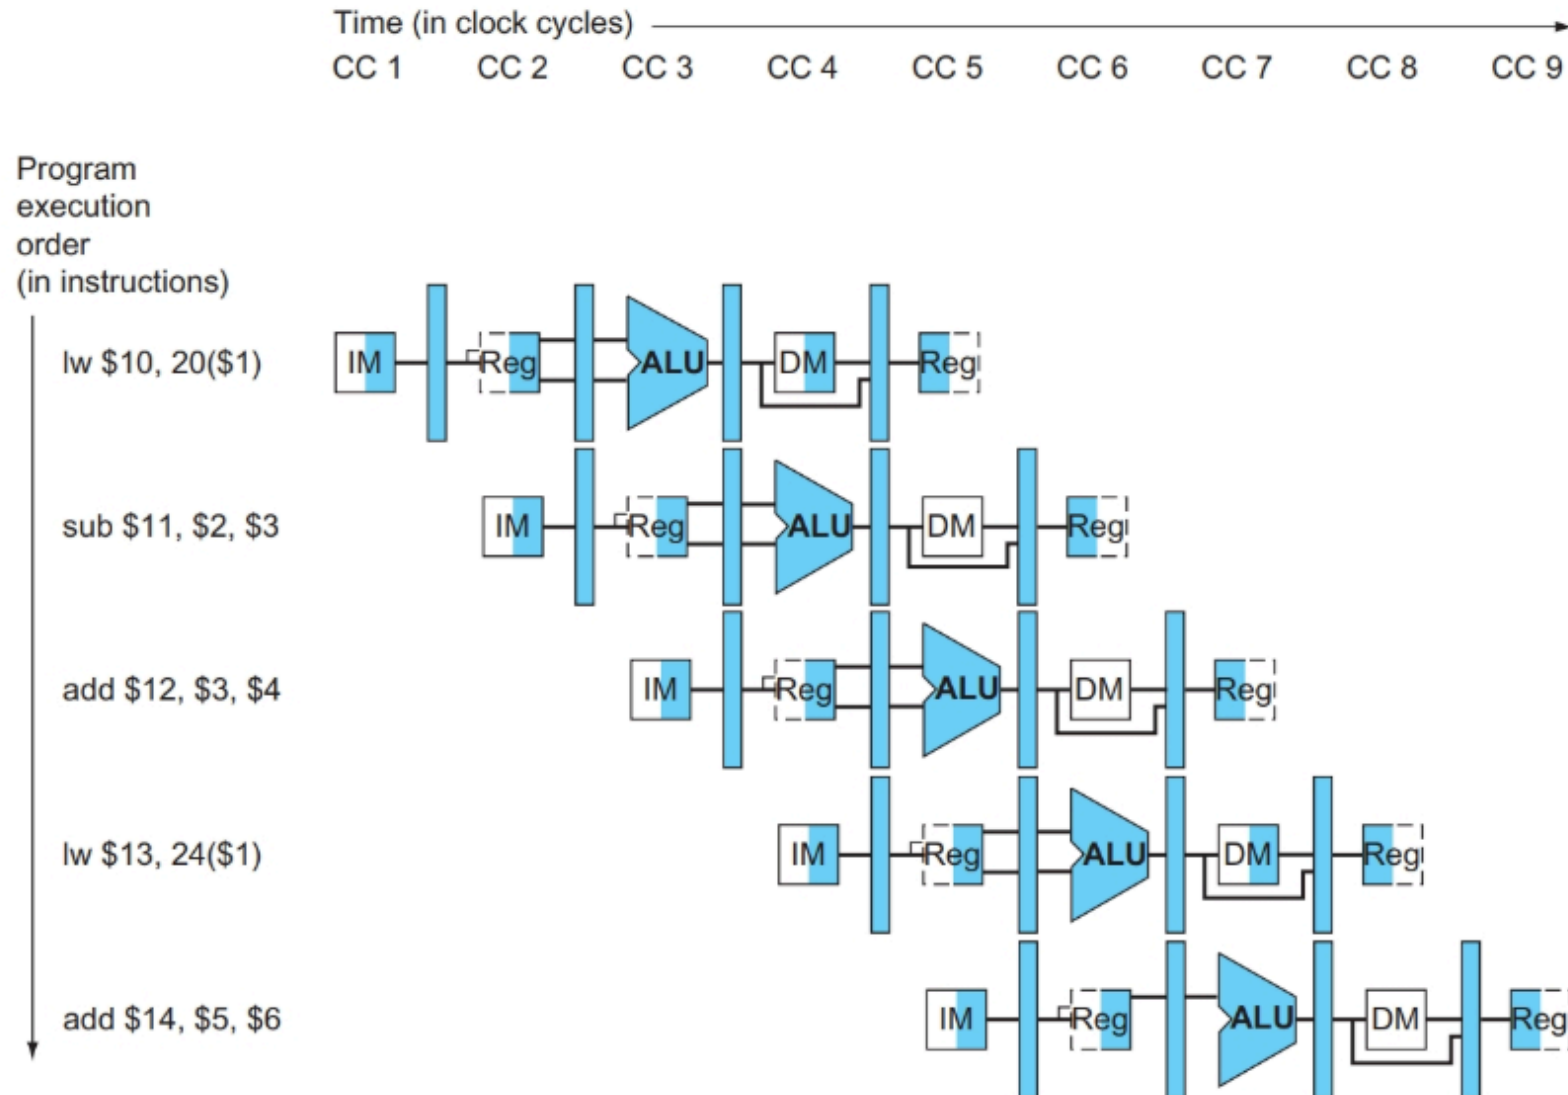
WB for Store



Example

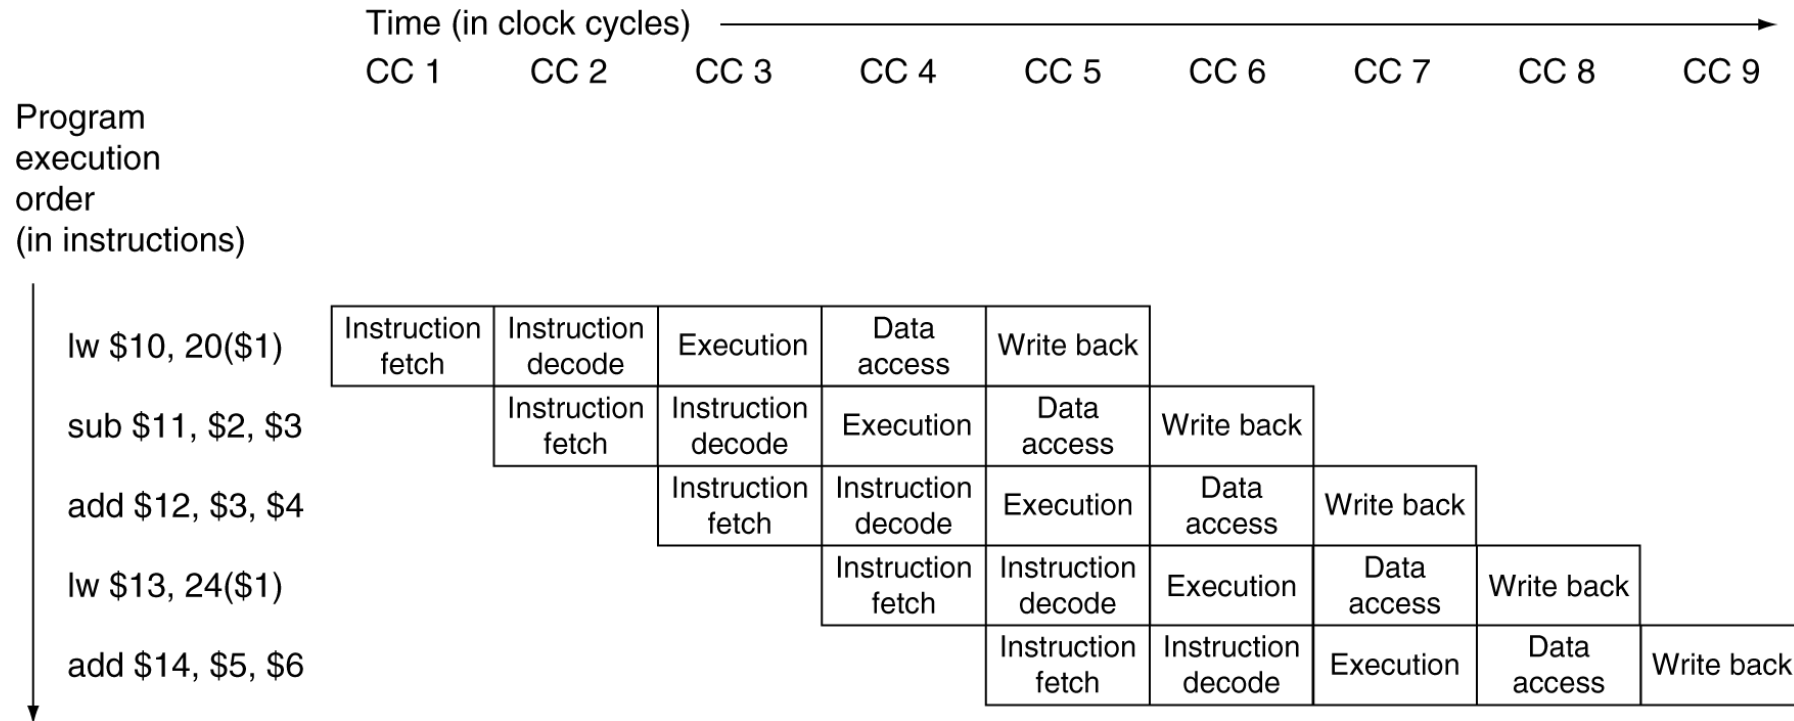
```
lw      $10, 20($1)
sub     $11, $2, $3
add     $12, $3, $4
lw      $13, 24($1)
add     $14, $5, $6
```

Multi-Cycle Pipeline Diagram



Multi-Cycle Pipeline Diagram

Traditional form




Single-Cycle Pipeline Diagram

State of pipeline in a given cycle (Clock cycle 5)

add \$14, \$5, \$6	lw \$13, 24(\$1)	add \$12, \$3, \$4	sub \$11, \$2, \$3	lw \$10, 20(\$1)
Instruction fetch	Instruction decode	Execution	Memory	Write-back

lw \$10, 20(\$1)
sub \$11, \$2, \$3
add \$12, \$3, \$4
lw \$13, 24(\$1)
add \$14, \$5, \$6






Control Signals

Instruction	Execution/address calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
	RegDst	ALUOp1	ALUOp0	ALUSrc	Branch	Mem- Read	Mem- Write	Reg- Write	Memto- Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
sw	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

Pipelined Control

- **Control signals derived from instruction**
 - As in single-cycle implementation



Pipelined Control

